

CSCE 435 Group project

0. Group number: 2 (Honors)

1. Group members:

1. Kimberly Chen
2. Spencer Le
3. Suhu Lavu
4. Andrew Mao
5. Jeff Ooi

We will communicate via a messaging group chat.

2. Project topic (e.g., parallel sorting algorithms)

Parallel Sorting Algorithms

2a. Brief project description (what algorithms will you be comparing and on what architectures)

We will be comparing the following algorithms:

- Bitonic Sort - Kimberly Chen
- Sample Sort - Spencer Le
- Merge Sort - Suhu Lavu
- Radix Sort - Andrew Mao
- Column Sort - Jeff Ooi

We will use the Grace cluster on the TAMU HPRC.

2b. Pseudocode for each parallel algorithm

- For MPI programs, include MPI calls you will use to coordinate between processes
- Bitonic Sort Pseudocode

```

func bitonic_sort(matrix, lowIndex, count, direction)
  Initialize MPI
  rank <- MPI rank
  num_procs <- MPI size

  if count > 1
    k = count / 2 // Divide the array into two halves

    // Sort the first half in ascending order
    bitonicSort(matrix, lowIndex, k, 1)

    // Sort the second half in descending order
    bitonicSort(matrix, lowIndex + k, k, 0)

    // Merge the two halves according to the 'direction'
    bitonicMerge(arr[], lowIndex, count, direction)

  for step = 1 to log2(num_procs)
    partner <- rank XOR step

    // Exchange sorted halves with partner process and merge
    if direction == 1 // this means it is ascending
      if rank < partner
        MPI send local_array to partner
        MPI receive partner_array from partner
        matrix <- bitonic_merge_ascending(matrix, partner_array)
      else
        MPI send local_array to partner
        MPI receive partner_array from partner
        matrix <- bitonic_merge_descending(matrix, partner_array)
    else // this means it is descending
      if rank < partner
        MPI send local_array to partner
        MPI receive partner_array from partner
        matrix <- bitonic_merge_descending(matrix, partner_array)
      else
        MPI send local_array to partner
        MPI receive partner_array from partner
        matrix <- bitonic_merge_ascending(matrix, partner_array)
    end if
  end for
end if

// Gather all sorted data at master

```

```

    if rank is master
        for i: 1 -> num_procs
            MPI receive sorted segments from all processes
        end for
        output "Final sorted array"
    end if

    Finalize MPI
end func

```

- Sample Sort Pseudocode

```

func parallel_sample_sort(local_data):
    // Initialize MPI environment
    MPI_Init()

    // Get total number of processes and rank of each process
    int num_procs, rank
    MPI_Comm_size(MPI_COMM_WORLD, &num_procs)
    MPI_Comm_rank(MPI_COMM_WORLD, &rank)

    // Step 1: Sort local data locally
    local_data = sort(local_data)

    // Step 2: Choose 'p' samples from the sorted local data
    samples = select_samples(local_data, num_procs)

    // Step 3: Gather all samples to the root process
    all_samples = []
    MPI_Gather(samples, num_procs, MPI_INT, all_samples, num_procs, MPI_INT, 0)

    // Step 4: Root process sorts the gathered samples and selects splitters
    splitters = []
    if rank == 0:
        all_samples_sorted = sort(all_samples)
        splitters = select_splitters(all_samples_sorted, num_procs)

    // Step 5: Broadcast splitters to all processes
    MPI_Bcast(splitters, num_procs, MPI_INT, root=0, MPI_COMM_WORLD)

    // Step 6: Each process redistributes its local data based on the splitters
    buckets = distribute_data(local_data, splitters, num_procs)

    // Step 7: Each process sends its buckets to the appropriate processes
    for i = 0 to num_procs - 1:
        // Send and receive buckets to/from each process
        send_bucket_to_process(buckets[i], i)
        receive_bucket_from_process(i)
    end for
end func

```

```

// Step 8: Merge the received buckets
local_data = merge(received_buckets)

// Step 9: Perform local sort on the merged data
local_data = sort(local_data)

// Step 10: Gather sorted data on the root process
sorted_data = []
MPI_Gather(local_data, len(local_data), MPI_INT, sorted_data, len(local_data), MPI_ROOT, MPI_COMM_WORLD)

// Finalize MPI environment
MPI_Finalize()

return sorted_data
end func

```

- Merge Sort Pseudocode

```

func parallel_merge_sort(array, n):
    Initialize MPI
    rank <- MPI rank
    size <- MPI size

    # calculate size of worker array and create local copy
    worker_size = n / size
    worker_arr = new array[worker_size]

    # send real array to all processes
    MPI_Scatter(array, worker_size, MPI_INT, worker_arr, worker_size, MPI_INT, rank, MPI_COMM_WORLD)

    # sort local copy
    sorted_arr = merge_sort(worker_arr)

    # merge sorted pieces back together
    step = 1
    while step < size:
        if rank % (2 * step) == 0:
            if rank + step < size:
                # get sorted array from another process
                received_size = worker_size * step
                received_array = new array[received_size]
                MPI_Recv(received_array, received_size, MPI_INT, source=rank + step, MPI_COMM_WORLD)

                # merge received array with local array
                sorted_arr = merge(sorted_arr, received_array)
                worker_size += received_size
            else:
                # no more data to receive, just sort the remaining data
                sorted_arr = merge_sort(sorted_arr)

        # send the sorted array back to the root process
        MPI_Send(sorted_arr, worker_size, MPI_INT, dest=0, MPI_COMM_WORLD)

        # double the step size
        step *= 2

    return sorted_arr

```

```

        else:
            # send sorted array to another process
            target = rank - step
            MPI_Send(sorted_arr, worker_size, MPI_INT, dest=target, tag=0, MPI_
            break out of loop
        step *= 2

# master process gathers the final sorted array
if rank == 0:
    MPI_Gather(sorted_arr, worker_size, MPI_INT, array, worker_size, MPI_INT

Finalize MPI
end func

func merge(arr1, arr2):
    size_1 = size(arr1)
    size_2 = size(arr2)
    final = new array[size_1 + size_2]
    i, j, k = 0

    while i < size_1 and j < size_2:
        if arr1[i] < arr2[j]:
            final[k] = arr1[i]
            i += 1
        else:
            final[k] = arr2[j]
            j += 1
        k += 1

    # copy leftover elements
    while i < size_1:
        final[k] = arr1[i]
        i += 1
        k += 1

    while j < size_2:
        final[k] = arr2[j]
        j += 1
        k += 1

    return final
end func

func merge_sort(arr):
    # base case
    if size(arr) <= 1:
        return arr

    # find midpoint

```

```

mid = size(arr) / 2

# sort left array
left_sorted = merge_sort(arr[0:mid])

# sort right array
right_sorted = merge_sort(arr[mid + 1:])

# merge them together
return merge(left_sorted, right_sorted)
end func

```

- Radix Sort Pseudocode

```

func radix_sort(matrix, lowIndex, count, direction)
  Initialize MPI
  rank <- MPI rank
  num_procs <- MPI size

  keysPerWorker <- numKeys / num_procs

  if rank is master:
    for i: 1 -> num_procs - 1:
      MPI send arr[i * keysPerWorker to (i + 1) * keysPerWorker - 1] to
    end for
  end if

  localbuffer <- [keysPerWorker]

  if rank is worker:
    MPI receive arr[i * keysPerWorker to (i + 1) * keysPerWorker - 1] from
  end if

  # For each iteration, process g bits at a time (Radix Sorting)
  max_bits <- get bits max(arr)
  for i: 0 -> (max_bits / g) do:
    local_bucket_counts <- counting_sort(localbuffer, g, i)

    # Aggregate the result of counting sort, make global count of keys per
    global_bucket_counts <- MPI all_to_all(local_bucket_counts)
    prefix_sums <- get prefix sums of global_bucket_counts
    bucketed_keys <- distribute_keys(localbuffer, g, i, prefix_sums)

    # update these results to preserve order
    exchanged_keys <- MPI all_to_all_exchange(bucketed_keys)
    sort_by_g_bits(exchanged_keys, g, i)
    localbuffer <- exchanged_keys
  end for
end func

```

```

end for

if rank is worker:
    MPI send localbuffer to master
end if

if rank is master:
    for i: 1 -> num_procs - 1 do:
        MPI receive sorted buffer into arr
    end for

    Finalize MPI
end func

```

- Column Sort Pseudocode

Assumption: $\text{numRows} \geq 2 * (\text{numCols} - 1)^2$

```
func column_sort(matrix, numRows, numCols)
```

```
    Initialize MPI
```

```
    rank <- MPI rank
```

```
    num_procs <- MPI size
```

```
    colsPerWorker = numCols / num_procs
```

```
    if rank is master
```

```
        for i: 1 -> num_procs
```

```
            MPI send matrix columns i * colsPerWorker to (i + 1) * colsPerWorker
```

```
        end for
```

```
    localbuffer <- buffer with enough size to fit matrix columns i * colsPerWorker
```

```
    if rank is not master
```

```
        MPI receive matrix columns i * colsPerWorker to (i + 1) * colsPerWorker
```

```
    sort every column in localbuffer
```

```
    if rank is not master
```

```
        MPI send sorted localbuffer to master
```

```
    if rank is master
```

```
        for i: 1 -> num_procs
```

```
            MPI receive all localbuffers into matrix
```

```
        end for
```

```
    transpose matrix
```

```
    for i: 1 -> num_procs
```

```
        MPI send matrix columns i * colsPerWorker to (i + 1) * colsPerWorker
    end for
```

```

if rank is not master
    MPI receive matrix columns i * colsPerWorker to (i + 1) * colsPerWorker

sort every column in localbuffer

if rank is not master
    MPI send sorted localbuffer to master
if rank is master
    for i: 1 -> num_procs
        MPI receive all localbuffers into matrix
    end for

    untranspose matrix

    for i: 1 -> num_procs
        MPI send matrix columns i * colsPerWorker to (i + 1) * colsPerWorker
    end for

if rank is not master
    MPI receive matrix columns i * colsPerWorker to (i + 1) * colsPerWorker

sort every column in localbuffer

if rank is not master
    MPI send sorted localbuffer to master
if rank is master
    for i: 1 -> num_procs
        MPI receive all localbuffers into matrix
    end for

    shift matrix

    for i: 1 -> num_procs
        MPI send matrix columns i * colsPerWorker to (i + 1) * colsPerWorker
    end for

if rank is not master
    MPI receive matrix columns i * colsPerWorker to (i + 1) * colsPerWorker

sort every column in localbuffer

if rank is not master
    MPI send sorted localbuffer to master
if rank is master
    sort last column in matrix

    unshift matrix // matrix should now be sorted in Column-Major order

```



```

        for i: 1 -> num_procs
            MPI send matrix columns i * colsPerWorker to (i + 1) * colsPerWorker
        end for
    if rank is not master
        MPI receive matrix columns i * colsPerWorker to (i + 1) * colsPerWorker

    check if localbuffer is sorted
    sorted <- 1 if sorted, 0 if not
    if rank is master
        isSorted <- -1
        MPI reduce sorted to master with minimum sorted into isSorted
        if isSorted is 1
            output "matrix sorted"
        else
            output "matrix not sorted"

    Finalize MPI
end func

```

2c. Evaluation plan - what and how will you measure and compare

We will measure and compare the sorting times of:

- Different input sizes and input types
 - Input sizes are 2^{16} , 2^{18} , 2^{20} , 2^{22} , 2^{24} , 2^{26} , and 2^{28}
 - Input types are sorted, random, reverse sorted, and 1% perturbed
- Strong scaling (same problem size, increase number of processors)
 - Number of processes are 2, 4, 8, 16, 32, 64, 128, 256, 512, and 1024
- Weak scaling (increase problem size, increase number of processors)
 - Number of processes are 2, 4, 8, 16, 32, 64, 128, 256, 512, and 1024

We will collect them using Caliper and compare them using Thicket.